

[15 hrs] BUILD: Practicum II / Mine a Database

In Part 1 you create a normalized relational OLTP database and populate it with data from an XML document. In Part 2 you will add to the normalized schema fact tables and turn the normalized schema into a de-normalized schema suitable for OLAP. In Part 3 you'll use the OLAP star/snowflake schema to do some (simple) data mining. The parenthesis contain the maximum number of points that can be earned for that question and an estimate of the time in hours.

Part 1 (55 pts) Load XML Data into Database

1. (0 pts / 0.1 hrs) Create an R Project named `CS5200.PracicumII.LastNameF` (where *LastName* is your last name and *F* is your first initial, e.g., `CS5200.PracicumII.WuX` or both last names plus initial, e.g., `CS5200.PracicumII.WuX-ChangY`, if you work in a pair) and then within that project create an R Script (.R file and not a .Rmd Notebook) called `LoadXML2DB.LastNameF.R` for Part 1 (same naming as above if you work in a group); Parts 2 and 3 will be done in different programs.
2. (0 pts / 0.5 hrs) Download the following XML file: [pubmed22n0001-tf.xml](#) [Links to an external site.](#). Save the XML files locally in a subfolder named `pubmed-tfm-xml` inside your project folder, and then inspect the file to familiarize yourself with its content and structure. You might wish to create a subset of the first file for use during development as it will load faster. To download and save an XML file, open the context menu by (generally) right-clicking on the link and selecting "Save As" or a similar choice depending on your operating system.
3. (5 pts / 45 min) Build an external DTD for the XML file above and reference that DTD from the XML downloaded in (2). Fewer points will be awarded if the DTD is internal.
4. (5 pts / 1 hr) Create a normalized relational schema that contains the following entities/tables: *Articles*, *Journals*, *Authors*. Use the XML document to determine the appropriate attributes (fields/columns) for the entities (tables). While there may be other types of publications in the XML, you only need to deal with articles in journals. Create appropriate primary and foreign keys. Where necessary, add synthetic surrogate keys. Store all information necessary for Parts 2 and 3. The choice of table structure is up to you. You may wish to visualize the schema using an ERD, although you do not have to submit the ERD.
5. (10 pts / 1 hr) Realize the relational schema in SQLite (use CREATE TABLE statements using R functions; you cannot use {sql} chunks in a R Script). You may **not use MySQL** for this step.
6. (5 pts / 2 hrs) Load (with validation of the DTD) the XML file into R from a URL (or locally; but fewer points will be awarded). You need to upload the XML and DTD to a hosting service of your choice.
7. (30 pts / 5 hrs) Extract and transform the data from the XML file in the folder and then save the data into the appropriate tables in the database. Given the structure of the XML, it is not feasible to only use `xmlToDataFrame` but instead will need parse the XML using a combination of node-by-node tree traversal and *XPath*. It is not feasible to use *XPath* to extract all journals, then all authors, etc. as some are missing and won't match up. You will need to iterate through the top-level nodes. While outside the scope of the course, this task could also be done through XSLT. Do not store duplicate authors or journals. For dates, you need to devise a conversion scheme, document your decision, and convert all dates to your chosen encoding scheme. You may wish to create a smaller XML file for testing rather than processing the larger XML file each time you run the code during development.
8. (5 pts / 0 hrs) Practice good coding: write R functions as needed; structure your code to be readable and maintainable; use proper programming practices, thoroughly comment your code; add your name, course information, date to a header in your script.

Part 2 (25 pts) Create Star/Snowflake Schema

1. (0 pts / 0.1 hrs) Create a new R Script for Part 2 named `LoadDataWarehouse.LastNameF.R` (use same naming as before if you are in a group).

2. (5 pts / 1 hr) Create a MySQL database using either a local or a cloud MySQL instance. Connect to the database. If you use SQLite for this step, no credit will be awarded for this question.
3. (20 pts / 5 hrs) Create and populate a star schema for journal facts. You may wish to visualize the star schema using an ERD, although you do not have to submit the ERD. Each row in this fact table will represent one journal fact. It must include the journal PK (as you've designed in Part 1 above), journal title, number of articles published in the journal per quarter and per year, and number of unique authors who published per quarter and per month. Load the data from the SQLite Database created in Part 1 and populate the fact table via R. Do not load the entire tables from SQLite into R as that does not scale. Use SQL to get the data you need. Note that there is not a single way to create the fact table -- you may use dimension tables or you may collapse the dimensions into the fact table. Remember that the goal of fact tables is to make interactive analytical queries fast through pre-computation and storage -- more storage but better performance. This requires thinking and creativity -- there is not a single best solution. Make sure your code scales -- imagine that your transactional database (the SQLite database) contains hundreds of millions of rows per table -- would your code still work and still performance well? ETL code is not always the fastest but it must scale. This [chalk talk](#)[Links to an external site.](#) on how to build fact tables and why might help...

Your star schema must support analytical queries such as these:

What are the number of articles published in every journal in 2012 and 2013?

What is the number of articles published in every journal in each quarter of 2012 through 2015?

How many articles were published each quarter (across all years)?

How many unique authors published articles in each year for which there is data?

This is the fact table you will need for Part 3.2. Of course, the years in the examples above are arbitrary and other analytical queries are possible, but these are examples to help you define your fact tables.

Part 3 (20 pts) Explore and Mine Data

1. (0 pts / 0.5 hrs) Create an R Notebook named *AnalyzeData.LastNameF.Rmd* within your R Project (use same naming as before if you are in a group).
2. (15 pts / 4 hrs) In the notebook, use markdown to write a "report" which shows the results of the following analytical queries against your MySQL data warehouse from Part 2 (which might go to some manager as part of a weekly report):

Analytical Query I: *Top five journals with the most articles published in them for the time period. Think about how you would best present this information so it is easily usable by an analyst or manager.*

Analytical Query II: *Number of articles per journal per year broken down by quarter. Again, think of a good way to show this. Is a table better than a visualization or is it better to use a visualization.*

Of course, be sure that your report uses the data from the fact tables.

If you use SQLite for this step, maximum half the credit will be awarded for this question. Each of the analytics use cases is equally weighted. Thoughtfulness and creativity do count and we want you to think about your presentation and report. Apply yourself a bit.